

 6 contributors

pipework / 

Download ZIP

- create a bridge named `br1` in the docker host;
- add an interface named `eth1` to the `$APACHE` container;
- assign IP address 192.168.1.1 to this interface,
- connect said interface to `br1`.

Now (drum roll), let's do this:

```
pipework br1 $MYSQL 192.168.1.2/24
```

This will:

- not create a bridge named `br1`, since it already exists;
- add an interface named `eth1` to the `$MYSQL` container;
- assign IP address 192.168.1.2 to this interface,
- connect said interface to `br1`.

Now, both containers can ping each other on the 192.168.1.0/24 subnet.

## Docker integration

Pipework can resolve Docker containers names. If the container ID that you gave to Pipework cannot be found, Pipework will try to resolve it with `docker inspect`. This makes it even simpler to use:

```
docker run -name web1 -d apache
pipework br1 web1 192.168.12.23/24
```

## Peeking inside the private network

Want to connect to those containers using their private addresses? Easy:

```
ip addr add 192.168.1.254/24 dev br1
```

Voilà!

## Setting container internal interface

By default pipework creates a new interface `eth1` inside the container. In case you want to change this interface name like `eth2`, e.g., to have more than one interface set by pipework, use:

```
pipework br1 -i eth2 ...
```

## Using a different netmask

The IP addresses given to `pipework` are directly passed to the `ip addr` tool; so you can append a subnet size using traditional CIDR notation.

I.e.:

```
pipework br1 $CONTAINERID 192.168.4.25/28
```

Don't forget that all containers should use the same subnet size; pipework is not clever enough to use your specified subnet size for the first container, and retain it to use it for the other containers.

## Setting a default gateway

If you want *outbound* traffic (i.e. when the containers connects to the outside world) to go through the interface managed by Pipework, you need to change the default route of the container.

This can be useful in some usecases, like traffic shaping, or if you want the container to use a specific outbound IP address.

This can be automated by Pipework, by adding the gateway address after the IP address and subnet mask:

```
pipework br1 $CONTAINERID 192.168.4.25/28@192.168.4.1
```

## Connect a container to a local physical interface

Let's pretend that you want to run two Hipache instances, listening on real interfaces eth2 and eth3, using specific (public) IP addresses. Easy!

```
pipework eth2 $(docker run -d hipache /usr/sbin/hipache) 50.19.169.157
pipework eth3 $(docker run -d hipache /usr/sbin/hipache) 107.22.148.5
```

Note that this will use `macvlan` subinterfaces, so you can actually put multiple containers on the same physical interface.

## Wait for the network to be ready

Sometimes, you want the extra network interface to be up and running *before* starting your service. A dirty (and unreliable) solution would be to add a `sleep` command before starting your service; but that could break in "interesting" ways if the server happens to be a bit slower at one point.

There is a better option: add the `pipework` script to your Docker image, and before starting the service, call `pipework --wait`. It will wait until the `eth1` interface is present and in `UP` operational state, then exit gracefully.

## Add the interface without an IP address

If for some reason you want to set the IP address from within the container, you can use `0/0` as the IP address. The interface will be created, connected to the network, and assigned to the container, but without configuring an IP address:

```
pipework br1 $CONTAINERID 0/0
```

## DHCP

You can use DHCP to obtain the IP address of the new interface. Just specify `dhcp` instead of an IP address; for instance:

```
pipework eth1 $CONTAINERID dhcp
```

You need three things for this to work correctly:

- obviously, a DHCP server (in the example above, a DHCP server should be listening on the network to which we are connected on `eth1`);
- the `udhcpd` DHCP client must be installed on your Docker *host* (you don't have to install it in your containers, but it must be present on the host);
- the underlying network must support bridged frames.

The last item might be particularly relevant if you are trying to bridge your containers with a WPA-protected WiFi network. I'm not 100% sure about this, but I think that the WiFi access point will drop frames originating from unknown MAC addresses; meaning that you have to go through extra hoops if you want it to work properly.

It works fine on plain old wired Ethernet, though.

## Specify a custom MAC address

If you need to specify the MAC address to be used (either by the `macvlan` subinterface, or the `veth` interface), no problem. Just add it as the command-line, as the last argument:

```
pipework eth0 $(docker run -d haproxy) 192.168.1.2 26:2e:71:98:60:8f
```

This can be useful if your network environment requires whitelisting your hardware addresses (some hosting providers do that), or if you want to obtain a specific address from your DHCP server. Also, some projects like [Orchestrator](#) rely on static MAC-IPv6 bindings for DHCPv6:

```
pipework br0 $(docker run -d zerorpcworker) dhcp fa:de:b0:99:52:1c
```

**Note:** if you generate your own MAC addresses, try remember those two simple rules:

- the lowest bit of the first byte should be `0`, otherwise, you are defining a multicast address;
- the second lowest bit of the first byte should be `1`, otherwise, you are using a globally unique (OUI enforced) address.

In other words, if your MAC address is `?X:??:??:??:??:??`, `X` should be `2`, `6`, `a`, or `e`. You can check [Wikipedia](#) if you want even more details.

## Support Open vSwitch

If you want to attach a container to the Open vSwitch bridge, no problem.

```
ovs-vsctl list-br
ovsbr0
pipework ovsbr0 $(docker run -d mysql /usr/sbin/mysqld_safe) 192.168.1.2/24
```

## Cleanup

When a container is terminated (the last process of the net namespace exits), the network interfaces are garbage collected. The interface in the container is automatically destroyed, and the interface in the docker host (part of the bridge) is then destroyed as well.

